

---

# Parallel Particle-in-Cell Algorithm

***CS 205 - Final Presentation***


**Kevin Howarth, Aditi Memani, Hari Raval, Taro Spirig**

---

# Motivation

# Background

- **Particle in Cell (PIC) is an attractive approach for simulating particle trajectories**
  - **N-body** →  **$O(N^2)$**  due to considering binary interactions between all particles
  - **PIC** →  **$O(N)$**  due to performing work on each particle only once
- **Princeton Plasma Physics Laboratory develops XGC**
  - PIC research code
  - Informs design of new fusion systems
  - Planned for use on exascale computing architectures



*A property heatmap inside a tokamak simulated by XGC*  
Source: [insidehpc.com](https://www.insidehpc.com)


# Scientific Goals

- **Demonstrate capability of our PIC code to simulate plasma behavior at realistically large input values**
  - Handle a large number of gridpoints → reduce spatial discretization error
  - Simulate at small timesteps → accurately resolve high frequency plasma behavior
- **Serve as a proof of concept for accurately representing plasma phenomena**
  - $10^{18}$  particles /  $\text{m}^3$  in real-world fusion problems

# Algorithm

# PIC Algorithm Overview


1. Interpolate from Particles to Mesh



2. Solve Discrete Poisson Equation on Mesh

$$\nabla^2 \phi = -\rho, \quad \nabla \phi = E$$

3. Interpolate from Mesh back to Particles



4. Time-step Particle Locations

$$\frac{d\bar{v}}{dt} = qE(\bar{x}), \quad \frac{d\bar{x}}{dt} = \bar{v}$$

# Data Structures

*Particle*: **Not** NUMA aware

```
Particle() {
```

```
    std::vector<float> xx;
```

```
    std::vector<float> xy;
```


```
    std::vector<float> vx;
```

```
    std::vector<float> vy;
```

```
    std::vector<float> Ex;
```

```
    std::vector<float> Ey;
```

```
}
```



```
Particle* Particles = new Particle[N_p]();
```

# Data Structures

*Particle: **Not** NUMA aware*

```
Particle() {
```

```
    std::vector<float> xx;  
    std::vector<float> xy;  
    std::vector<float> vx;  
    std::vector<float> vy;  
    std::vector<float> Ex;  
    std::vector<float> Ey;  
}
```

*Particle: **NUMA** aware*

```
Particle(int N_tin) {
```

```
    Nt = N_tin;  
    xx = new float [Nt];  
    xy = new float [Nt];  
    vx = new float [Nt];  
    vy = new float [Nt];  
    Ex = new float [Nt];  
    Ey = new float [Nt];  
}
```




*Particle\* Particles = new Particle[N\_p]();*



# Data Structures

## Computational Mesh




<b>Charge</b>	<b>Potential</b>	<b>Electric Field</b>	$\in \mathbb{R}^{N_g \times N_g}$
---------------	------------------	-----------------------	-----------------------------------

*Eigen:: Matrix qGrid, pGrid, eGrid;*

# Parallel Implementation

# Step 1: Particle-to-Mesh Interpolation

## Algo. Steps



## Parallel Implementation of P2M

- P2M reads from particles and writes to (shared) mesh
  - Implement thread-private grids and custom reduction
  - Parallelize over number of particles

$$\phi_{M4}(x) = \begin{cases} 1 - \frac{5}{2}|x|^2 + \frac{3}{2}|x|^3 & \text{if } |x| \leq 1 \\ \frac{1}{2}(2 - |x|)^2(1 - |x|) & \text{if } 1 < |x| \leq 2 \\ 0 & 2 \leq |x| \end{cases}$$


***M4 Interpolation Kernel***

$$\rho_d = q_p \times \phi_{M4}(x_{p,d}) \times \phi_{M4}(y_{p,d})$$

***Rule to update grid point “d” with signature from particle “p”***

# Step 1: Particle-to-Mesh Interpolation

## Algo. Steps



## Parallel Implementation of P2M

- P2M reads from particles and writes to (shared) mesh
  - Implement thread-private grids and custom reduction
  - Parallelize over number of particles


```
/*-----*/  
// declare a custom reduction which performs addition of Eigen::Matrix objects  
#pragma omp declare reduction(MatrixPlus: Eigen::Matrix: omp_out=omp_out+omp_in) \  
| initializer(omp_priv=Eigen::Matrix::Zero(omp_orig.rows(),omp_orig.cols()))  
/*-----*/
```

```
// perform the custom reduction to add Eigen::Matrix objects  
#pragma omp parallel for reduction(MatrixPlus: chargeGrid)  
  for (unsigned int p = 0; p < N_p; p++) {  
    P2M_M4_kernel(particles[p], chargeGrid);  
  }  
}
```

*Custom reduction definition and usage*

# Step 2: Successive Over-Relaxation

## Algo. Steps




## Parallel Implementation of SOR

- SOR corresponds to stencil:
  - Read from Self and Neighbors, Write to Self
- Parallelization over number of grid points in Red-Black sweeps

$$V_{i,j}^{k+1} = (1 - \omega^k)V_{i,j}^k + \omega^k \frac{V_{i-1,j}^k + V_{i+1,j}^k + V_{i,j-1}^k + V_{i,j+1}^k + h^2 \rho_{i,j}}{4}$$


*SOR update rule for grid point (i, j) in iteration "k"*

*Example stencil for update during black sweep*



# Step 2: Successive Over-Relaxation


## Algo. Steps



## Parallel Implementation of SOR

```
// update "BLACK" nodes in "red-black" ordering scheme
#pragma omp parallel for schedule(static)
    for (int i = 0; i < n; i++) {
        for (int j = 0; j < n - 1; j += 2) {
            V[i][j] = ... // SOR Stencil Here
        }
    }
```


*Parallelization of black sweep*



*Visualization of parallel black sweep + importance of memory layout!*

# Step 3: Mesh-to-Particle Interpolation

## Algo. Steps



## Parallel Design of M2P

- M2P reads from shared mesh and writes to particles
- Parallelize over the number of particles

$$\vec{E}_p = \sum_{d \in \mathbb{D}_{M4}} \vec{E}_d \times \phi_{M4}(x_{p,d}) \times \phi_{M4}(y_{p,d})$$


*Rule to update particle “p” with signature from surrounding grid points  $d \in \mathbb{D}_{M4}$*

```
#pragma omp parallel for schedule(static)
    for (unsigned int p = 0; p < N_p; p++) {
        |
        |     M2P_M4_kernel(particles[p], electricFieldGrid);
        |
    }
}
```

*Parallelization of M2P interpolation*

# Step 4: Time-stepping

## Algo. Steps



## Parallel Design of Time-stepping

- Time-step using Leapfrog Method
  - Read from and write to same particle
  - Independent per particle → parallelize over number of particles


$$\begin{aligned}\vec{v}_{t+1/2} &= \vec{v}_{t-1/2} + \Delta t \frac{q}{m} \vec{E}(\vec{x}_t) \\ \vec{x}_{t+1} &= \vec{x}_t + \Delta t \vec{v}_{t+1/2}\end{aligned}$$

***Leapfrog Method rule to update a given particle's kinematic data***



# Step 4: Time-stepping

## Algo. Steps



## Parallel Design of Time-stepping

- Possible load imbalance when enforcing periodic wraparounds
- Address via dynamic scheduling with large chunk size

```
#pragma omp parallel for schedule(dynamic, 500)
for (unsigned int p = 0; p < N_p; p++) {
    // timestep the velocity
    particles[p].vx[ti + 1] = particles[p].vx[ti] + constantFactorV * particles[p].Ex[ti];
    particles[p].vy[ti + 1] = particles[p].vy[ti] + constantFactorV * particles[p].Ey[ti];

    // timestep the position
    float xNew = particles[p].x[ti] + constantFactorX * particles[p].vx[ti + 1];
    float yNew = particles[p].y[ti] + constantFactorX * particles[p].vy[ti + 1];

    // enforce periodic boundary condition
    while (particleOutsideGrid) {
        periodicWraparound_x(xNew, particleOutsideGrid);
        periodicWraparound_y(yNew, particleOutsideGrid);
    }

    // write the new particle position
    particles[p].x[ti + 1] = xNew;
    particles[p].y[ti + 1] = yNew;
}
```

# Parallel Results

# Problem Size

## *Small Case*

$10^6$  particles ( $N_p$ )  
400 grid points ( $N_g$ )  
100 timesteps ( $N_t$ )

## *Large Case*

$3 \times 10^7$  particles ( $N_p$ )  
1500 grid points ( $N_g$ )  
100 timesteps ( $N_t$ )

## *Memory Requirement*

$$\text{Mem}(N_p, N_g, N_t) = 4 \times (N_p \times N_t \times (2 + 2 + 2) + N_g^2 \times (1 + 1 + 2)) \times 10^{-9}$$

***Note: All performance benchmarks and scaling analysis run on 1 Intel Broadwell node (~126 GB memory) with the large problem size***

# Problem Size

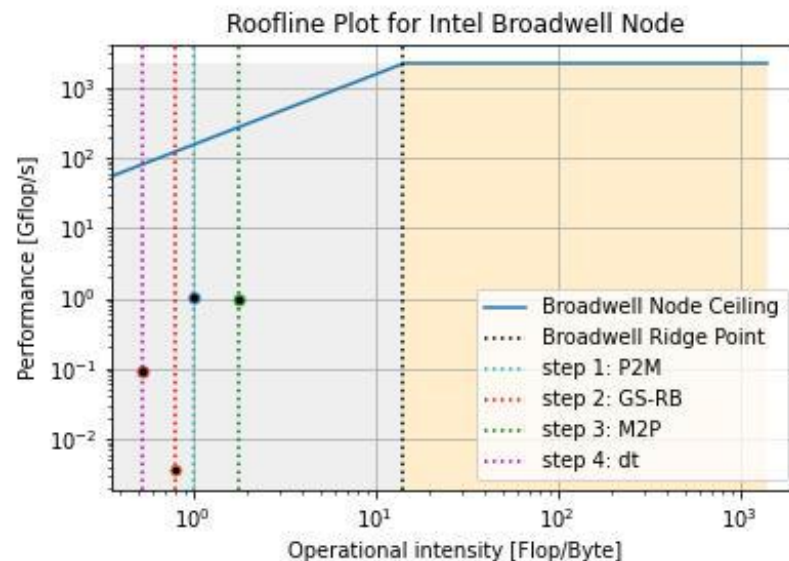
<b><i>Problem Size</i></b>	<b>"Small" Case</b>	<b>"Large" Case</b>
<b><i>Wall clock time [hours]</i></b>	0.03	1.5
<b><i>Memory per node [GB]</i></b>	2.4	72.1

***Our "large" problem size requires 72.1 GB which is less than the Intel Broadwell node capacity, and is still computationally reasonable***

# Roofline Analysis: Sequential

- All four kernels are memory bound
- Gauss-Seidel - largest bottleneck with performance  $< 0.01$  Gflop/s
- Focus on analyzing *M2P*, because it has highest ceiling for improvement

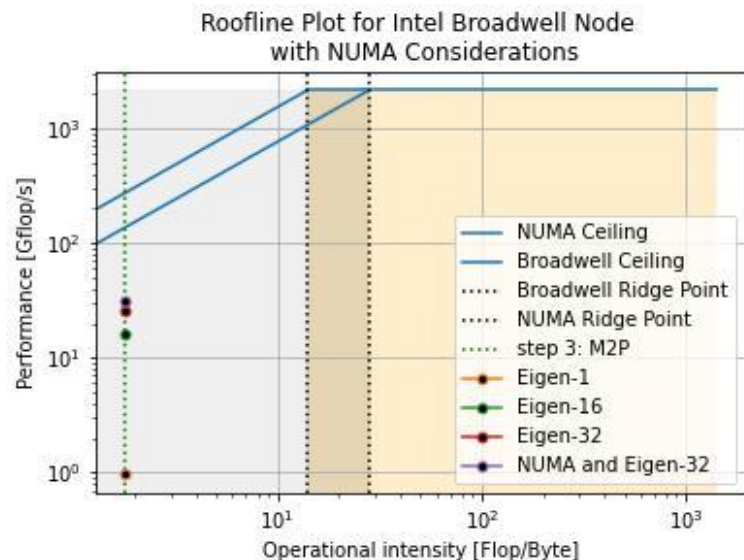
## Roofline for Sequential Code



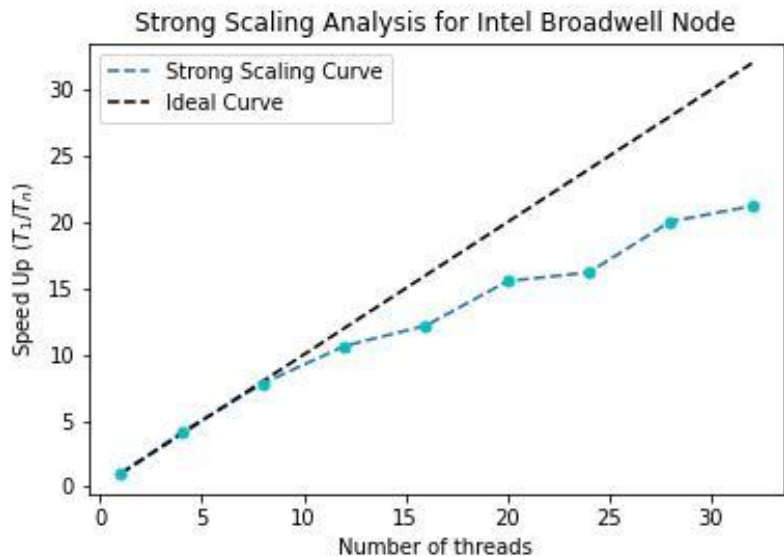
# Roofline Analysis: M2P

- Improves from 1.0 Gflop/s (sequentially) to 32.7 Gflop/s (32 threads, NUMA)
- Difference between NUMA aware and unaware is small
  - M2P infrequently interacts with the *Particle* data structure
- Other kernels (e.g. Timestepping) with more *Particle* memory accesses benefit significantly from NUMA

## Roofline for Parallel M2P Code



# Strong Scaling Analysis



Threads	Time ( $T_n$ )	Speed Up ( $T_1/T_n$ )
1	5214.2856 s.	1.0
4	1260.3978 s.	4.1
8	665.9816 s.	7.8
12	489.5262 s.	10.7
16	427.8973 s.	12.2
20	335.5052 s.	15.5
24	321.7918 s.	16.2
28	260.1357 s.	20.0
32	245.8093 s.	21.2

**Implementing shared memory parallelism with OpenMP, leveraging the Eigen library, and including NUMA first touch considerations produced a speed-up of ~21.2X with 32 threads**

# Conclusions



# Key Takeaways

- **Implement and optimize high-performance code using:**
  - **Tools:** OpenMP, Eigen, Pybind11
  - **Techniques:** Data layout/access patterns, scheduling, effective data structures
  - **Software-Hardware Interplay:** Roofline Analysis, NUMA first-touch policy
- **Kevin will conduct his master's thesis implementing a new style of PIC algorithm with a group at the Princeton Plasma Physics Laboratory**

# Future Work: Holistic Improvements

- **Implement efficient I/O pipeline**
  - Increase usability and efficiency
  - Explore capabilities available with HDF5 library
- **Switch to SoA data structures to enable SIMD vectorization**
- **Migrate to MPI + OpenMP Hybrid memory model**
  - Enables much larger problem sizes

# Future Work: Kernel-Specific Improvements

- **Particle to Mesh Interpolation:**
  - Leverage partially shared grids or cell list data structures (*reduce memory footprint*)
  - Implement a higher order interpolation kernel (*increase accuracy and operational intensity*)
- **Discrete Poisson Equation:**
  - Use a better algorithm such as multigrid or FFT (*increase accuracy and efficiency*)
- **Mesh to Particle Interpolation:**
  - Implement a higher order interpolation kernel (*increase accuracy and operational intensity*)
- **Timestepping:**
  - Utilize a higher order ODE Solver (*increase accuracy and operational intensity*)

**Thank You**